

# KGCS: Building a Frozen Knowledge Graph for Expert Cybersecurity AI

**Version:** 1.0

**Date:** January 2026

**Purpose:** Design, architecture, and implementation of a standards-backed cybersecurity knowledge graph that enables safe, explainable, and hallucination-free AI reasoning about vulnerabilities, attacks, defenses, and threat intelligence.

---

## Executive Summary

**KGCS** (Cybersecurity Knowledge Graph) is a **frozen, immutable, standards-aligned ontology** that integrates 9 MITRE security taxonomies (CVE, CWE, CPE, CVSS, CAPEC, ATT&CK, D3FEND, CAR, SHIELD, ENGAGE) to create a single source of truth for cybersecurity AI systems.

## Why This Matters

Current cybersecurity AI systems suffer from:

- **Semantic drift:** Standards are reinterpreted to fit models, losing original meaning
- **Hallucination:** LLMs invent causal links that don't exist in source data
- **Opacity:** Answers can't be traced back to authoritative sources
- **Fragility:** New standard releases break downstream systems
- **Scope creep:** Adding one feature (e.g., incidents) requires redesigning the entire graph

**KGCS solves this** by enforcing three invariants:

1. **Authoritative alignment:** Every ontology class maps 1:1 to official JSON/STIX schemas
2. **Explicit provenance:** Every relationship is traceable to source data
3. **Layered architecture:** Core (frozen facts) + Extensions (contextual, temporal, inferred)

The result is an AI that can reason confidently, explain its reasoning, and remain maintainable as standards evolve.

---

## Table of Contents

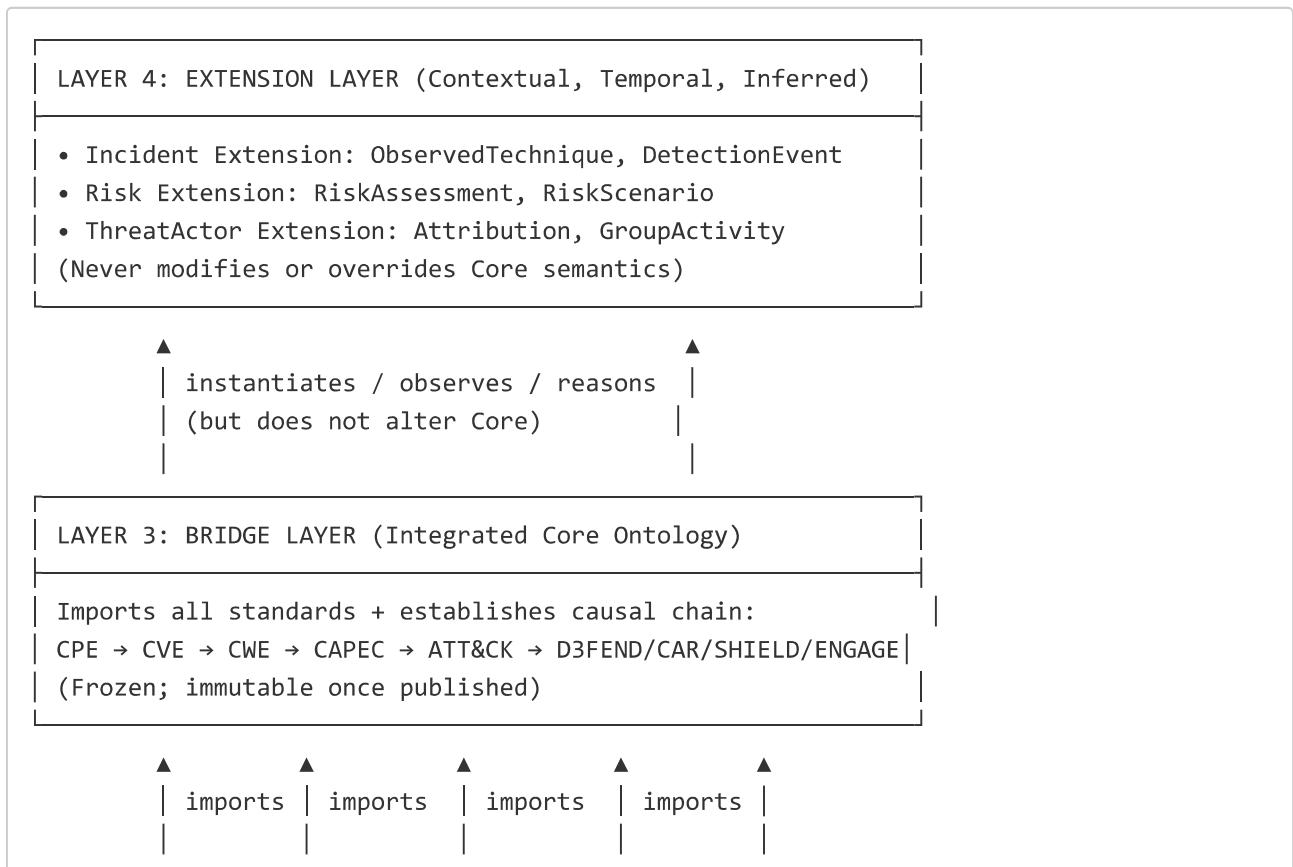
- 1. Core Architecture
- 2. The Vulnerability Causality Chain

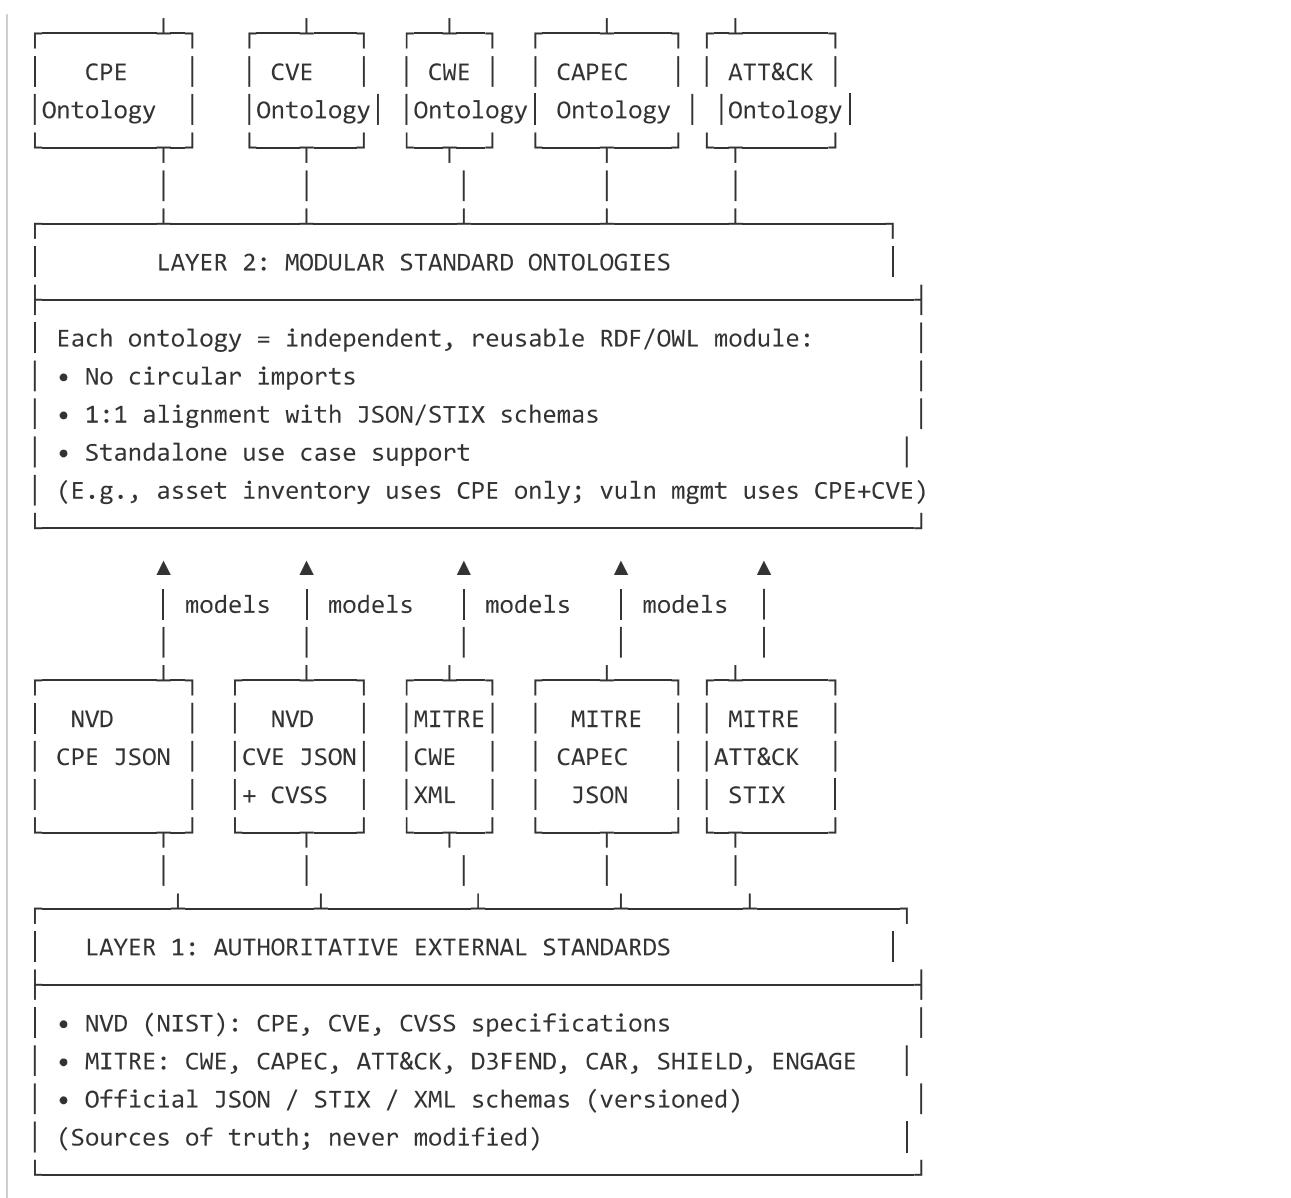
- 3. Defense, Detection, and Deception Coverage
  - 4. Engagement and Strategic Reasoning
  - 5. Standards Alignment by Layer
  - 6. Modular Ontology Design
  - 7. RAG-Safe Traversal Paths
  - 8. Extension Layers (Incident, Risk, Threat Actor)
  - 9. Implementation Roadmap
  - 10. Why This Architecture Scales
  - 11. Security Guarantees
  - 12. Limitations and Future Work
  - 13. For the AI Engineer
  - Conclusion
- 

# 1. Core Architecture

## 1.1 Layered Ontology Design

KGCS is built on **four immutable layers**, each with a clear semantic boundary:





## 1.2 Core Principle: External Standards = Authoritative Truth

**Rule 1:** Every ontology class maps 1:1 to a JSON object or field group in official standards.

**Rule 2:** Every relationship is explicitly traceable to source data.

**Rule 3:** No invented semantics. The ontology is a **lens**, not a replacement.

## 2. The Vulnerability Causality Chain

The **canonical path** through the knowledge graph is:

```

CPE (Asset Identifier)
  ↓ affected_by
CVE (Vulnerability Disclosure)
  └── scored_by → CVSS (Severity Assessment)
    └── caused_by
CWE (Weakness Root Cause)
  ↓ exploited_by
CAPEC (Attack Pattern Abstraction)
  
```

```

└ implements_as / maps_to
ATT&CK Technique (Adversary Behavior)
└ mitigated_by → D3FEND (Defensive Technique)
└ detected_by → CAR (Detection Analytic)
└ countered_by → SHIELD (Deception Technique)
└ disrupted_by → ENGAGE (Strategic Concept)

```

## 2.1 Why This Chain Matters

This path **enforces explainability**. Every step:

- Exists in authoritative JSON/STIX
- Has a documented source
- Is independently verifiable
- Cannot be skipped without losing information

**Example Question:** "How does CVE-2025-1234 relate to ATT&CK Technique T1059?"

**Answer (with provenance):**

```

CVE-2025-1234
→ caused_by: CWE-94 (Improper Control of Generation of Code)
  [source: cvejson.weaknesses[]]
→ exploited_by: CAPEC-242 (Code Injection)
  [source: capec JSON Related_Weaknesses]
→ maps_to: ATT&CK T1059 (Command and Scripting Interpreter)
  [source: MITRE-documented mapping]

```

**Every hop is verifiable.**

## 2.2 Critical Invariant: Vulnerabilities Affect Configurations, Not Platforms

**Not this:**

```
CVE → Platform (wrong)
```

**But this:**

```

CVE → PlatformConfiguration
  (version bounds, update level, criteria)
  ↓
  includes → Platform

```

**Why?** A platform (e.g., "Windows 10") is not vulnerable. A *specific configuration* of that platform is (e.g., "Windows 10 version 21H2 without KB5012234").

NVD encodes this as logical expressions in `configurations[].nodes[].cpeMatch[]`. We preserve that.

## 3. Defense, Detection, and Deception Coverage

### 3.1 Parallel Control Dimensions

Three independent dimensions answer three questions about ATT&CK techniques:

Dimension	Question	Standard	Semantics
Mitigation	How to prevent or reduce impact?	D3FEND	Defensive technique
Detection	How to detect execution?	CAR	Detection analytic + data sources
Deception	How to deceive or manipulate?	SHIELD	Deception technique

These are **not hierarchical**. All three can apply to the same technique.

### 3.2 Example: T1087 (Account Discovery)

```
T1087 (ATT&CK)
└ mitigated_by → D3-ID-001 (Credential Access Restriction)
└ mitigated_by → D3-OT-002 (Harden User Account Privilege)
└ detected_by → CAR-2016-12-001 (Create Remote Process)
└ detected_by → CAR-2013-10-002 (SMB Write Request)
└ countered_by → SHIELD-002 (Fake Account Indicators)
```

**No coverage is perfect.** This graph lets an SOC:

- See all available defenses at a glance
- Choose which to implement based on cost/effectiveness
- Detect when gaps exist
- Reason about layered defense

## 4. Engagement and Strategic Reasoning

### 4.1 ENGAGE Does Not Map to Techniques

Unlike D3FEND, CAR, and SHIELD, ENGAGE operates at a **strategic level**, not a technical one.

```
ENGAGE Concept
└ targets → Threat Actor Group
└ informs → Response Strategy
└ influences → when/how D3FEND/CAR/SHIELD are deployed
```

**ENGAGE answers:**

- How do we interact with adversaries?
- What are their decision points?
- How do we disrupt their operations at a macro level?

**4.2 Example: Attribution as Engagement**

```
ENGAGE Strategy: "Coordinated Public Attribution"
└ targets: Nation-state actors
└ objective: Increase operational costs
└ informs: Decision to deploy SHIELD (deception)
    to gather intelligence for attribution
```

**5. Standards Alignment by Layer****5.1 Layer 1: CPE (Platform Identification)****Source:** NVD CPE Dictionary JSON 2.3**Ontology Class:** Platform (CPE)**Key Properties:**

- cpeUri (canonical string)
- part, vendor, product, version, update, edition, ...
- deprecated (boolean)

**Edges:**

- Platform —deprecates—> Platform (obsolescence tracking)
- Platform —isVariantOf—> Platform (version relationships)

**Use Case:** "What software/hardware is in our infrastructure?"**5.2 Layer 2: CVE + CVSS (Vulnerability Disclosure)****Source:** NVD CVE API 2.0 + CVSS JSON**Ontology Classes:**

- Vulnerability (CVE)
- VulnerabilityScore (CVSS) (versioned: v2, v3.1, v4.0)

- Reference (provenance)

### Key Properties:

- cveId, description, published, lastModified
- vulnStatus (Analyzed, Undergoing Analysis, Deferred, Disputed, Rejected)
- cisaDateAdded, cisaActionDue (federal remediation mandates)

### Critical:

- Each CVSS version = **separate node** (never overwrite)
- affected\_by relationship goes to PlatformConfiguration, not Platform

**Use Case:** "Is our system vulnerable to CVE-2025-1234?"

---

## 5.3 Layer 3: CWE (Weakness Root Cause)

**Source:** MITRE CWE JSON

**Ontology Class:** Weakness (CWE)

### Key Properties:

- cweId, name, abstraction (Pillar, Class, Base, Variant)
- description, status

### Edges:

- Weakness —parent\_of—> Weakness (hierarchical)
- Weakness —member\_of—> WeaknessView
- Weakness —exploited\_by—> AttackPattern

**Use Case:** "What underlying flaw does this CVE exploit?"

---

## 5.4 Layer 4: CAPEC (Attack Pattern Abstraction)

**Source:** MITRE CAPEC JSON

**Ontology Class:** AttackPattern (CAPEC)

### Key Properties:

- capecId, name, description
- likelihood, severity, prerequisites

**Edges:**

- `AttackPattern —exploits—> Weakness` (weakness → behavior bridge)
- `AttackPattern —related_to—> AttackPattern`
- `AttackPattern —maps_to—> Technique` (to ATT&CK)

**Use Case:** "How are weaknesses exploited in practice?"

---

**5.5 Layer 5: ATT&CK (Adversary Tradecraft)**

**Source:** MITRE ATT&CK STIX 2.1

**Ontology Classes:**

- `Technique` , `SubTechnique` (attack-pattern type)
- `Tactic` (x-mitre-tactic)
- `Group` , `Software` (actor / campaign context)
- `DataSource` , `DataComponent` (detection metadata)

**Key Relationships:**

- `Technique —part_of—> Tactic`
- `SubTechnique —subtechnique_of—> Technique` (functional)
- `Group —uses—> Technique`
- `Technique —detected_by—> DataComponent`

**Use Case:** "What are adversaries actually doing?"

---

**5.6 Layer 6: D3FEND (Defensive Techniques)**

**Source:** MITRE D3FEND STIX

**Ontology Class:** `DefensiveTechnique` (D3FEND)

**Subtypes:**

- `DetectionTechnique` (how to detect)
- `DenialTechnique` (how to block)
- `DisruptionTechnique` (how to interfere)

**Key Properties:**

- `d3fendId`, `name`, `sophisticationLevel`, `costLevel`, `scope`
- `implementationStatus` (Proposed, Beta, Stable, Deprecated)

**Edges:**

- `DefensiveTechnique` —mitigates—> `Technique` (direct)
- `DefensiveTechnique` —mitigates—> `Weakness` (root cause)
- `DefensiveTechnique` —weakens—> `AttackPattern`

**Use Case:** "How do we defend against this technique?"

---

**5.7 Layer 7: CAR (Detection Analytics)**

**Source:** MITRE CAR JSON/YAML

**Ontology Class:** `DetectionAnalytic` (CAR)

**Key Properties:**

- `carId`, `name`, `description`
- `techniques[]` (which ATT&CK techniques it detects)
- `data_sources[]` (what data it needs)

**Edges:**

- `DetectionAnalytic` —detects—> `Technique`
- `DetectionAnalytic` —requires—> `DataSource`

**Use Case:** "How do we detect this behavior?"

---

**5.8 Layer 8: SHIELD (Deception Techniques)**

**Source:** MITRE SHIELD STIX (experimental)

**Ontology Class:** `DeceptionTechnique` (SHIELD)

**Subtypes:**

- `HoneypotDeception` (attract and deceive)
- `MisdirectionDeception` (misdirect)
- `InformationDecoy` (false data)
- `SocialManipulation` (social engineering)

**Key Properties:**

- `shieldId`, `name`, `targetAdversary`
- `sophisticationLevel`, `deploymentComplexity`
- `primaryObjective` (Detect, Disrupt, Deceive, Gather Intelligence)

**Edges:**

- `DeceptionTechnique` —counters—> `Technique`
- `DeceptionTechnique` —reveals—> `Technique` (expose TTPs)

**Use Case:** "How do we deceive or observe adversaries?"

---

## 5.9 Layer 9: ENGAGE (Strategic Engagement)

**Source:** MITRE ENGAGE Framework (conceptual)

**Ontology Class:** `EngagementConcept` (ENGAGE)

**Subtypes:**

- `DisruptionStrategy` (disrupt operations)
- `OperationalInterference` (interfere with C2)
- `CapabilityDegradation` (weaken capabilities)
- `AttributionStrategy` (identify and attribute)
- `ResilienceBuilding` (strengthen defenses)

**Key Properties:**

- `engageId`, `name`, `strategyType`
- `timeframe` (Immediate, Short-term, Medium-term, Long-term)
- `operationalLevel` (Tactical, Operational, Strategic)
- `riskLevel`, `legalConsiderations`

**Edges:**

- `EngagementConcept` —disrupts—> `Technique`
- `EngagementConcept` —targets—> `Group`
- `EngagementConcept` —informs—> `ResponseStrategy`

**Use Case:** "How do we interact strategically with adversaries?"

## 6. Modular Ontology Design

### 6.1 Why Separate Ontologies?

Instead of one monolithic ontology, KGCS uses **modular, independent OWL files**:

```
docs/ontology/owl/
├── cpe-ontology-v1.0.owl      (standalone)
├── cve-ontology-v1.0.owl      (imports CPE)
├── cwe-ontology-v1.0.owl      (standalone)
├── capec-ontology-v1.0.owl    (imports CWE)
├── attck-ontology-v1.0.owl    (standalone)
├── d3fend-ontology-v1.0.owl   (imports Core)
├── car-ontology-v1.0.owl     (imports Core)
├── shield-ontology-v1.0.owl   (imports Core)
├── engage-ontology-v1.0.owl   (imports Core)
└── core-ontology-extended-v1.0.owl (imports all above)
  └── [extensions]
    ├── incident-ontology-extension-v1.0.owl
    ├── risk-ontology-extension-v1.0.owl
    └── threatactor-ontology-extension-v1.0.owl
```

### 6.2 Benefits of Modularity

Benefit	Enables
<b>Reusability</b>	Asset inventory systems use CPE alone; vulnerability management uses CPE+CVE
<b>Independent Versioning</b>	CPE 3.0 can be added without breaking CVE/CWE/etc.
<b>Scope Limiting</b>	Teams work on one standard without touching others
<b>Testing</b>	Each ontology validated independently before integration
<b>Future Growth</b>	New standards (CVSS-NG, D3FEND extensions) slot in cleanly

### 6.3 Import Graph (No Circular Dependencies)

```
Core Ontology
├── imports: cpe-ontology-v1.0.owl
├── imports: cve-ontology-v1.0.owl (imports CPE)
├── imports: cwe-ontology-v1.0.owl
├── imports: capec-ontology-v1.0.owl (imports CWE)
├── imports: attck-ontology-v1.0.owl
├── imports: d3fend-ontology-v1.0.owl
├── imports: car-ontology-v1.0.owl
└── imports: shield-ontology-v1.0.owl
```

```

└ imports: engage-ontology-v1.0.owl

Extensions (never in Core; always reference Core)
├ incident-ontology-extension-v1.0.owl (imports Core)
├ risk-ontology-extension-v1.0.owl (imports Core)
└ threatactor-ontology-extension-v1.0.owl (imports Core)

```

**Critical Rule:** Extensions **reference** Core, never the reverse.

---

## 7. RAG-Safe Traversal Paths

### 7.1 Pre-Approved Query Templates

To prevent RAG hallucination, KGCS defines **approved traversal templates**. These enforce:

1. No layer skipping
2. No circular reasoning without context
3. Only authoritative edges

**Example Template T-CORE-01:** "How does this CVE impact this platform?"

```

Query: CVE-X → Platform-Y?
├ CVE-X —affects—> PlatformConfiguration-Z
|   (source: cvejson.configurations[].cpeMatch[].vulnerable)
└ PlatformConfiguration-Z —includes—> Platform-Y
|   (source: cpematch_api_json_2.0.schema)
└ Result: Yes, if configuration matches

```

**Example Template T-CORE-02:** "What are all defenses against this technique?"

```

Query: Defenses for Technique-T?
├ Technique-T —mitigated_by—> DefensiveTechnique
|   (source: MITRE D3FEND stix relationships)
└ Technique-T —detected_by—> DetectionAnalytic
|   (source: CAR techniques[])
└ Technique-T —countered_by—> DeceptionTechnique
|   (source: SHIELD mappings)

```

### 7.2 Forbidden Traversals

These are explicitly disallowed:

Forbidden	Why
CVE → ATT&CK (direct)	Must pass through CWE→CAPEC
Platform → CWE (direct)	Must pass through CVE

Forbidden	Why
Weakness → Group (direct)	No direct relationship; contextual only
Technique → Risk Score (direct)	Risk is extension-layer only

---

## 8. Extension Layers (Incident, Risk, Threat Actor)

### 8.1 Why Extensions Exist

Core ontology captures **what is known to be true** (facts from standards).

Extensions capture **what we observe or infer** (facts from operations/analysis).

**Critical:** Extensions **never modify Core classes**. They reference them.

### 8.2 Incident Extension

**Purpose:** Track temporal, contextual observations from SIEM/SOAR

**Key Classes:**

- `ObservedTechnique` (timestamp, confidence, evidence)
- `DetectionEvent` (alert, rule, sensor)
- `IncidentTimeline` (temporal sequence)

**Relationship to Core:**

```
ObservedTechnique —instantiates—> Technique (ATT&CK)
DetectionEvent —detects—> ObservedTechnique
IncidentTimeline —includes—> DetectionEvent
```

**Properties:**

- `timestamp`, `firstSeen`, `lastSeen`, `confidence` (LOW|MEDIUM|HIGH)
- `sourceSystem`, `evidenceIds[]`

### 8.3 Risk Extension

**Purpose:** Capture risk assessment, prioritization, and decision-making

**Key Classes:**

- `RiskAssessment` (scenario + score)
- `RiskScenario` (vulnerability → impact)

- `RemediationDecision` (ACCEPT|MITIGATE|TRANSFER|AVOID)

### Relationship to Core:

```
RiskScenario
└─ involves→ Vulnerability (CVE)
└─ affects→ Asset
└─ exploits→ AttackPattern (CAPEC)
└─ mitigated_by→ DefensiveTechnique (D3FEND)
```

### Properties:

- `riskScore` (0-100), `likelihood`, `impact`
- `decision`, `decisionRationale`, `decisionDate`
- `owner`, `owner_contact`

## 8.4 ThreatActor Extension

**Purpose:** Track attribution claims with confidence levels

### Key Classes:

- `AttributionClaim` (claim about group responsibility)
- `ThreatActorObservation` (observed behavior)

### Relationship to Core:

```
AttributionClaim
└─ attributes_to→ Group (ATT&CK)
└─ based_on→ ObservedTechnique
└─ confidence→ (LOW|MEDIUM|HIGH|VERY_HIGH)
└─ evidenceIds[]→ ThreatActorObservation
```

### Properties:

- `claimId`, `timestamp`, `confidence`
- `sourceSystem` (intel feed, SOAR, analyst), `analyst_note`

## 9. Implementation Roadmap

### Phase 1: Core Standards (Complete ✓)

- CPE Ontology (v1.0)
- CVE Ontology (v1.0)

- CWE Ontology (v1.0)
- CAPEC Ontology (v1.0)
- ATT&CK Ontology (v1.0)
- D3FEND Ontology (v1.0)
- CAR Ontology (v1.0)
- SHIELD Ontology (v1.0)
- ENGAGE Ontology (v1.0)
- Core Ontology Extended (v1.0, imports all above)

**Deliverable:** 10 OWL/Turtle files with full 1:1 alignment to JSON/STIX schemas.

## Phase 2: SHACL Validation (Next)

- Create SHACL shapes for each ontology
- Define constraint profiles: SOC (observational), EXEC (decision-focused), AI (hallucination-safe)
- Integrate pre-ingest validation (reject non-conforming data)
- Add runtime validation gates

**Deliverable:** SHACL shape files + validation pipeline

## Phase 3: Data Ingestion (Planned)

- ETL scripts to load NVD CVE/CPE JSON
- MITRE STIX 2.1 parser for ATT&CK / D3FEND / SHIELD
- CWE / CAPEC / CAR JSON loaders
- Neo4j or RDF graph database setup
- Re-ingestion safety tests (versioning, conflict detection)

**Deliverable:** Production-ready data pipeline

## Phase 4: Extension Layers (Planned)

- Incident Ontology Extension + temporal reasoning
- Risk Ontology Extension + prioritization models
- ThreatActor Ontology Extension + attribution logic

- RAG Traversal Template library

**Deliverable:** Safe, extensible query framework

## Phase 5: AI Integration (Planned)

- RAG retrieval layer (safe traversals only)
- Explanation generation (path-based)
- Confidence scoring (traceability to source)
- Fine-tuning LLMs on KGCS queries

**Deliverable:** Hallucination-free cybersecurity AI

---

# 10. Why This Architecture Scales

## 10.1 Handling Standards Evolution

**Scenario:** CVSS 5.0 is released tomorrow.

**Old approach:** Rewrite vulnerability model, re-ingest all CVEs, hope nothing breaks.

**KGCS approach:**

1. Add `CVSSv50Score` subclass to `VulnerabilityScore` in CVE ontology
2. Load new NVD data with v5.0 metrics
3. Existing CVSS v3.1 / v4.0 nodes remain unchanged
4. RAG queries automatically see both versions

**Why this works:** Different CVSS versions are separate nodes with `scored_by` edges.

## 10.2 Adding New Standards

**Scenario:** You want to add NIST SP 800-53 (security controls).

**KGCS approach:**

1. Create `controls-ontology-v1.0.owl` (new, standalone)
2. Define edges: `DefensiveTechnique —implements—> Control`
3. Import into Core Ontology
4. Existing queries unaffected; new queries enabled

**Why this works:** Modular design + no inheritance = clean extension.

## 10.3 Supporting Organizational Context

**Scenario:** You have 10,000 assets to integrate.

**KGCS approach:**

1. Assets live in **Asset Extension**, not Core
2. Core remains pristine (no CMDB data)
3. Asset —configured\_with—> PlatformConfiguration (from Core)
4. Asset —affected\_by—> Vulnerability (transitive)
5. Organizational changes don't break the graph

**Why this works:** Clear separation of authoritative (Core) vs. contextual (Extension).

## 10.4 Versioning and Rollback

**Scenario:** You ingest bad data into CAR, want to rollback one version.

**KGCS approach:**

1. Every entity has `cveId`, `capecId`, `carId` (stable external ID)
2. Nodes are immutable; only edges change
3. Version control at the RDF level:

```
# Before: CAR-2013-10-002 —detects—> T1007 (with 3 evidence trails)
# Problem found: 1 trail is wrong
# After: Remove bad edge, keep 2 valid edges
# Rollback: Replay transaction, done
```

**Why this works:** Every relationship is traceable to source + timestamped.

---

## 11. Security Guarantees

### 11.1 No Hallucination

Every answer is a **path through nodes and edges**, each backed by:

- Stable external ID (CVE-YYYY-XXXX, T1234, etc.)
- Source document (JSON field, STIX property)
- Timestamp (when data was ingested)

### 11.2 Explainability

Every query result includes:

```
{
  "answer": "CVE-2025-1234 relates to T1059",
  "confidence": "HIGH (authoritative path)",
  "path": [
    {
      "node": "CVE-2025-1234",
      "source": "NVD JSON id field"
    },
    {
      "edge": "caused_by",
      "source": "NVD JSON weaknesses[].description[]"
    },
    ...
  ],
  "evidence_urls": ["https://nvd.nist.gov/vuln/detail/CVE-2025-1234", ...]
}
```

## 11.3 Auditability

- Every node change is timestamped
- Every edge change is logged with provenance
- Data can be replayed from standard sources
- No hidden inferences

## 11.4 Compliance

- CVSS, CWE, CAPEC, ATT&CK mappings certified by MITRE
- CPE, CVE, CVSS alignment certified by NVD
- No custom interpretations
- Suitable for:
  - Regulatory reporting (PCI-DSS, NIST CSF, ISO 27001)
  - SOC playbooks and runbooks
  - Executive decision-making

# 12. Limitations and Future Work

## 12.1 Current Scope

KGCS 1.0 covers:

- Core standards (CVE, CWE, CPE, CVSS, CAPEC, ATT&CK, D3FEND, CAR, SHIELD, ENGAGE)
- 1:1 alignment with JSON/STIX schemas
- Modular OWL ontologies
- Formal semantics + constraints
- SHACL validation (in progress)
- Data ingestion pipeline (planned)
- RAG integration (planned)

## 12.2 Intentional Exclusions

These are **contextual, not authoritative**, so they belong in Extensions:

- Threat actor motivations (organizational intel)
- Incident timelines (operational data)
- Asset inventory (CMDB)
- Risk scores (business decision)
- Detection rules (SOC implementation)
- Exploit code (ethical concerns)

## 12.3 Future Enhancements

1. **Formal Verification:** Prove consistency of ontology using theorem provers
2. **Temporal Reasoning:** Track when relationships were discovered/confirmed
3. **Uncertainty Quantification:** Bayesian inference over relationship weights
4. **Fine-Grained Provenance:** Link individual CVSS metrics to scoring rationale
5. **Multi-Language Support:** Translate ontology labels/descriptions to non-English
6. **Linked Data Publishing:** Expose as SPARQL endpoints for external systems

# 13. For the AI Engineer

## 13.1 Using KGCS for RAG

```
# Query: "What techniques can compromise this system?"
```

1. Extract asset from user context  
Asset → PlatformConfiguration → CVE
2. For each CVE:  
CVE → CWE → CAPEC → ATT&CK Technique

3. For each Technique:
  - Technique → D3FEND (defenses)
  - Technique → CAR (detection)
  - Technique → SHIELD (deception)
  
4. Filter by:
  - Applicable to this environment (platform)
  - Confidence level (HIGH for authoritative paths)
  - Available data sources (CAR)
  
5. Return with full provenance:
 

```
[  
  {  
    "technique": "T1234",  
    "path": [...],  
    "defenses": [...],  
    "detection": [...],  
    "confidence": "HIGH"  
  }  
]
```

## 13.2 Fine-Tuning LLMs

Use KGCS to:

1. **Ground training data** in authoritative sources
  
  2. **Filter hallucinations** (reject answers not in path)
  
  3. **Score confidence** (authoritative path = HIGH)
  
  4. **Generate explanations** (path = explanation)
- 

## Conclusion

**KGCS is built for one purpose: Enable AI systems to reason confidently about cybersecurity.**

By aligning ontologies 1:1 with standards, enforcing explicit provenance, and layering contextual extensions cleanly, KGCS ensures that:

- Every answer is traceable to source
  
- Standards evolve without breaking the graph
  
- New contexts (assets, incidents, risks) integrate cleanly
  
- AI can explain its reasoning with evidence
  
- Humans retain control and auditability

**The result is an expert cybersecurity AI that is trustworthy, maintainable, and compliant.**

**For more information:**

- See [docs/ontology/](#) for formal ontology specifications
- See [.github/copilot-instructions.md](#) for AI agent governance
- See [docs/draft/](#) for detailed design documents