

KGCS: Building a Frozen Knowledge Graph for Expert Cybersecurity AI

Version: 1.0

Date: January 2026

Purpose: Design, architecture, and implementation of a standards-backed cybersecurity knowledge graph that enables safe, explainable, and hallucination-free AI reasoning about vulnerabilities, attacks, defenses, and threat intelligence.

Executive Summary

KGCS (Cybersecurity Knowledge Graph) is a **frozen, immutable, standards-aligned ontology** that integrates 9 MITRE security taxonomies (CVE, CWE, CPE, CVSS, CAPEC, ATT&CK, D3FEND, CAR, SHIELD, ENGAGE) to create a single source of truth for cybersecurity AI systems.

Why This Matters

Current cybersecurity AI systems suffer from:

- **Semantic drift:** Standards are reinterpreted to fit models, losing original meaning
- **Hallucination:** LLMs invent causal links that don't exist in source data
- **Opacity:** Answers can't be traced back to authoritative sources
- **Fragility:** New standard releases break downstream systems
- **Scope creep:** Adding one feature (e.g., incidents) requires redesigning the entire graph

KGCS solves this by enforcing three invariants:

1. **Authoritative alignment:** Every ontology class maps 1:1 to official JSON/STIX schemas
2. **Explicit provenance:** Every relationship is traceable to source data
3. **Layered architecture:** Core (frozen facts) + Extensions (contextual, temporal, inferred)

The result is an AI that can reason confidently, explain its reasoning, and remain maintainable as standards evolve.

Table of Contents

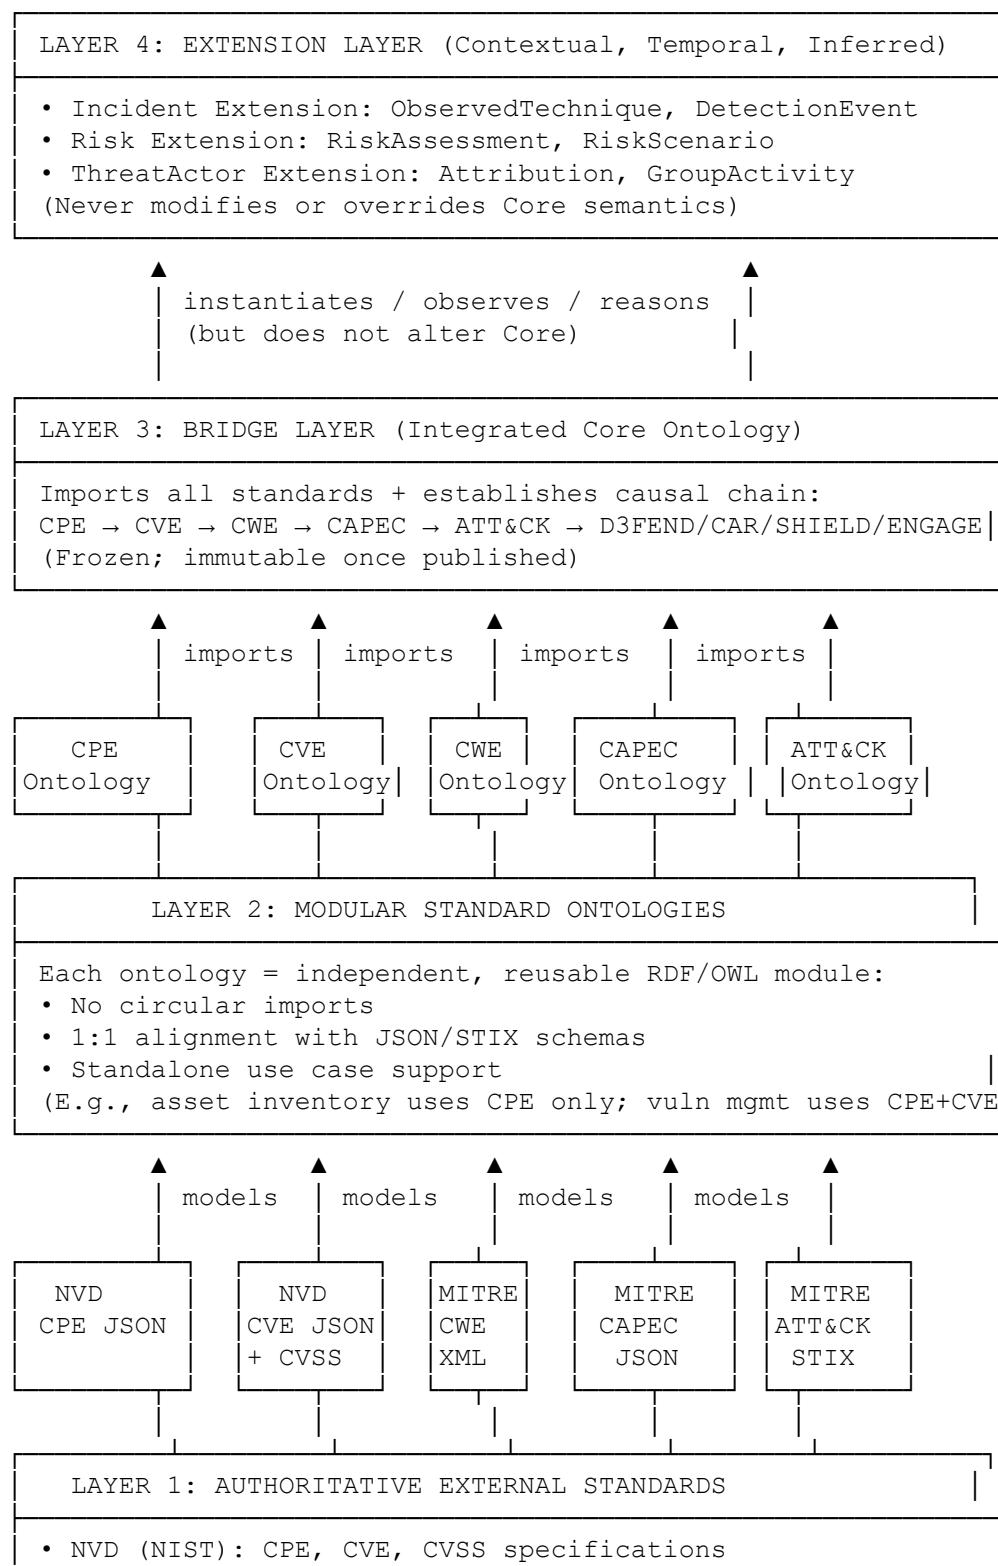
- [1. Core Architecture](#)
- [2. The Vulnerability Causality Chain](#)
- [3. Defense, Detection, and Deception Coverage](#)
- [4. Engagement and Strategic Reasoning](#)
- [5. Standards Alignment by Layer](#)
- [6. Modular Ontology Design](#)
- [7. RAG-Safe Traversal Paths](#)
- [8. Extension Layers \(Incident, Risk, Threat Actor\)](#)
- [9. Implementation Roadmap](#)
-

- [10. Why This Architecture Scales](#)
 - [11. Security Guarantees](#)
 - [12. Limitations and Future Work](#)
 - [13. For the AI Engineer](#)
 - [Conclusion](#)
-

1. Core Architecture

1.1 Layered Ontology Design

KGCS is built on **four immutable layers**, each with a clear semantic boundary:



- MITRE: CWE, CAPEC, ATT&CK, D3FEND, CAR, SHIELD, ENGAGE
- Official JSON / STIX / XML schemas (versioned)
- (Sources of truth; never modified)

1.2 Core Principle: External Standards = Authoritative Truth

Rule 1: Every ontology class maps 1:1 to a JSON object or field group in official standards.

Rule 2: Every relationship is explicitly traceable to source data.

Rule 3: No invented semantics. The ontology is a **lens**, not a replacement.

2. The Vulnerability Causality Chain

The **canonical path** through the knowledge graph is:

```

CPE (Asset Identifier)
  ↓ affected_by
CVE (Vulnerability Disclosure)
  ┌── scored_by → CVSS (Severity Assessment)
  └── caused_by
CWE (Weakness Root Cause)
  ↓ exploited_by
CAPEC (Attack Pattern Abstraction)
  ┌── implements_as / maps_to
ATT&CK Technique (Adversary Behavior)
  ├── mitigated_by → D3FEND (Defensive Technique)
  ├── detected_by → CAR (Detection Analytic)
  └── countered_by → SHIELD (Deception Technique)
  └── disrupted_by → ENGAGE (Strategic Concept)

```

2.1 Why This Chain Matters

This path **enforces explainability**. Every step:

- Exists in authoritative JSON/STIX
- Has a documented source
- Is independently verifiable
- Cannot be skipped without losing information

Example Question: "How does CVE-2025-1234 relate to ATT&CK Technique T1059?"

Answer (with provenance):

```

CVE-2025-1234
  → caused_by: CWE-94 (Improper Control of Generation of Code)
    [source: cvejson.weaknesses[]]
  → exploited_by: CAPEC-242 (Code Injection)
    [source: capec JSON Related_Weaknesses]
  → maps_to: ATT&CK T1059 (Command and Scripting Interpreter)
    [source: MITRE-documented mapping]

```

Every hop is verifiable.

2.2 Critical Invariant: Vulnerabilities Affect Configurations, Not Platforms

Not this:

CVE → Platform (wrong)

But this:

```
CVE → PlatformConfiguration  
    (version bounds, update level, criteria)  
    ↓  
    includes → Platform
```

Why? A platform (e.g., "Windows 10") is not vulnerable. A *specific configuration* of that platform is (e.g., "Windows 10 version 21H2 without KB5012234").

NVD encodes this as logical expressions in `configurations[].nodes[].cpeMatch[]`. We preserve that.

3. Defense, Detection, and Deception Coverage

3.1 Parallel Control Dimensions

Three independent dimensions answer three questions about ATT&CK techniques:

Dimension	Question	Standard	Semantics
Mitigation	How to prevent or reduce impact?	D3FEND	Defensive technique
Detection	How to detect execution?	CAR	Detection analytic + data sources
Deception	How to deceive or manipulate?	SHIELD	Deception technique

These are **not hierarchical**. All three can apply to the same technique.

3.2 Example: T1087 (Account Discovery)

```
T1087 (ATT&CK)  
└ mitigated_by → D3-ID-001 (Credential Access Restriction)  
└ mitigated_by → D3-OT-002 (Harden User Account Privilege)  
└ detected_by → CAR-2016-12-001 (Create Remote Process)  
└ detected_by → CAR-2013-10-002 (SMB Write Request)  
└ countered_by → SHIELD-002 (Fake Account Indicators)
```

No coverage is perfect. This graph lets an SOC:

- See all available defenses at a glance
- Choose which to implement based on cost/effectiveness
- Detect when gaps exist
- Reason about layered defense

4. Engagement and Strategic Reasoning

4.1 ENGAGE Does Not Map to Techniques

Unlike D3FEND, CAR, and SHIELD, ENGAGE operates at a **strategic level**, not a technical one.

ENGAGE Concept

- └ targets → Threat Actor Group
- └ informs → Response Strategy
- └ influences → when/how D3FEND/CAR/SHIELD are deployed

ENGAGE answers:

- How do we interact with adversaries?
- What are their decision points?
- How do we disrupt their operations at a macro level?

4.2 Example: Attribution as Engagement

ENGAGE Strategy: "Coordinated Public Attribution"

- └ targets: Nation-state actors
- └ objective: Increase operational costs
- └ informs: Decision to deploy SHIELD (deception) to gather intelligence for attribution

5. Standards Alignment by Layer

5.1 Layer 1: CPE (Platform Identification)

Source: NVD CPE Dictionary JSON 2.3

Ontology Class: Platform (CPE)

Key Properties:

- cpeUri (canonical string)
- part, vendor, product, version, update, edition, ...
- deprecated (boolean)

Edges:

- Platform —deprecates—> Platform (obsolescence tracking)
- Platform —isVariantOf—> Platform (version relationships)

Use Case: "What software/hardware is in our infrastructure?"

5.2 Layer 2: CVE + CVSS (Vulnerability Disclosure)

Source: NVD CVE API 2.0 + CVSS JSON

Ontology Classes:

- Vulnerability (CVE)
- VulnerabilityScore (CVSS) (versioned: v2, v3.1, v4.0)
- Reference (provenance)

Key Properties:

- cveId, description, published, lastModified
- vulnStatus (Analyzed, Undergoing Analysis, Deferred, Disputed, Rejected)
- cisaDateAdded, cisaActionDue (federal remediation mandates)

Critical:

- Each CVSS version = **separate node** (never overwrite)
- affected_by relationship goes to PlatformConfiguration, not Platform

Use Case: "Is our system vulnerable to CVE-2025-1234?"

5.3 Layer 3: CWE (Weakness Root Cause)

Source: MITRE CWE JSON

Ontology Class: Weakness (CWE)

Key Properties:

- cweId, name, abstraction (Pillar, Class, Base, Variant)
- description, status

Edges:

- Weakness —parent_of—> Weakness (hierarchical)
- Weakness —member_of—> WeaknessView
- Weakness —exploited_by—> AttackPattern

Use Case: "What underlying flaw does this CVE exploit?"

5.4 Layer 4: CAPEC (Attack Pattern Abstraction)

Source: MITRE CAPEC JSON

Ontology Class: AttackPattern (CAPEC)

Key Properties:

- capecId, name, description
- likelihood, severity, prerequisites

Edges:

- AttackPattern —exploits—> Weakness (weakness → behavior bridge)
- AttackPattern —related_to—> AttackPattern
- AttackPattern —maps_to—> Technique (to ATT&CK)

Use Case: "How are weaknesses exploited in practice?"

5.5 Layer 5: ATT&CK (Adversary Tradecraft)

Source: MITRE ATT&CK STIX 2.1

Ontology Classes:

- Technique, SubTechnique (attack-pattern type)
- Tactic (x-mitre-tactic)
- Group, Software (actor / campaign context)
- DataSource, DataComponent (detection metadata)

Key Relationships:

- Technique —part_of—> Tactic
- SubTechnique —subtechnique_of—> Technique (functional)
- Group —uses—> Technique
- Technique —detected_by—> DataComponent

Use Case: "What are adversaries actually doing?"

5.6 Layer 6: D3FEND (Defensive Techniques)

Source: MITRE D3FEND STIX

Ontology Class: DefensiveTechnique (D3FEND)

Subtypes:

- DetectionTechnique (how to detect)
- DenialTechnique (how to block)
- DisruptionTechnique (how to interfere)

Key Properties:

- d3fendId, name, sophisticationLevel, costLevel, scope
- implementationStatus (Proposed, Beta, Stable, Deprecated)

Edges:

- DefensiveTechnique —mitigates—> Technique (direct)
- DefensiveTechnique —mitigates—> Weakness (root cause)
- DefensiveTechnique —weakens—> AttackPattern

Use Case: "How do we defend against this technique?"

5.7 Layer 7: CAR (Detection Analytics)

Source: MITRE CAR JSON/YAML

Ontology Class: DetectionAnalytic (CAR)

Key Properties:

- carId, name, description
- techniques [] (which ATT&CK techniques it detects)
- data_sources [] (what data it needs)

Edges:

- DetectionAnalytic —detects—> Technique
- DetectionAnalytic —requires—> DataSource

Use Case: "How do we detect this behavior?"

5.8 Layer 8: SHIELD (Deception Techniques)

Source: MITRE SHIELD STIX (experimental)

Ontology Class: DeceptionTechnique (SHIELD)

Subtypes:

- HoneypotDeception (attract and deceive)
- MispirectionDeception (misdirect)
- InformationDecoy (false data)
- SocialManipulation (social engineering)

Key Properties:

- shieldId, name, targetAdversary
- sophisticationLevel, deploymentComplexity
- primaryObjective (Detect, Disrupt, Deceive, Gather Intelligence)

Edges:

- DeceptionTechnique —counters—> Technique
- DeceptionTechnique —reveals—> Technique (expose TTPs)

Use Case: "How do we deceive or observe adversaries?"

5.9 Layer 9: ENGAGE (Strategic Engagement)

Source: MITRE ENGAGE Framework (conceptual)

Ontology Class: EngagementConcept (ENGAGE)

Subtypes:

- DisruptionStrategy (disrupt operations)
- OperationalInterference (interfere with C2)
- CapabilityDegradation (weaken capabilities)
- AttributionStrategy (identify and attribute)
- ResilienceBuilding (strengthen defenses)

Key Properties:

- engageId, name, strategyType
- timeframe (Immediate, Short-term, Medium-term, Long-term)
- operationalLevel (Tactical, Operational, Strategic)
- riskLevel, legalConsiderations

Edges:

- EngagementConcept —disrupts—> Technique
- EngagementConcept —targets—> Group
- EngagementConcept —informs—> ResponseStrategy

Use Case: "How do we interact strategically with adversaries?"

6. Modular Ontology Design

6.1 Why Separate Ontologies?

Instead of one monolithic ontology, KGCS uses **modular, independent OWL files**:

```
docs/ontology/owl/
├── cpe-ontology-v1.0.owl          (standalone)
├── cve-ontology-v1.0.owl          (imports CPE)
├── cwe-ontology-v1.0.owl          (standalone)
├── capec-ontology-v1.0.owl        (imports CWE)
├── attck-ontology-v1.0.owl        (standalone)
├── d3fend-ontology-v1.0.owl       (imports Core)
├── car-ontology-v1.0.owl          (imports Core)
├── shield-ontology-v1.0.owl       (imports Core)
├── engage-ontology-v1.0.owl        (imports Core)
└── core-ontology-extended-v1.0.owl (imports all above)
  [extensions]
    ├── incident-ontology-extension-v1.0.owl
    ├── risk-ontology-extension-v1.0.owl
    └── threatactor-ontology-extension-v1.0.owl
```

6.2 Benefits of Modularity

Benefit	Enables
Reusability	Asset inventory systems use CPE alone; vulnerability management uses CPE+CVE
Independent Versioning	CPE 3.0 can be added without breaking CVE/CWE/etc.
Scope Limiting	Teams work on one standard without touching others
Testing	Each ontology validated independently before integration
Future Growth	New standards (CVSS-NG, D3FEND extensions) slot in cleanly

6.3 Import Graph (No Circular Dependencies)

```
Core Ontology
└── imports: cpe-ontology-v1.0.owl
└── imports: cve-ontology-v1.0.owl (imports CPE)
└── imports: cwe-ontology-v1.0.owl
└── imports: capec-ontology-v1.0.owl (imports CWE)
└── imports: attck-ontology-v1.0.owl
└── imports: d3fend-ontology-v1.0.owl
└── imports: car-ontology-v1.0.owl
└── imports: shield-ontology-v1.0.owl
└── imports: engage-ontology-v1.0.owl

Extensions (never in Core; always reference Core)
└── incident-ontology-extension-v1.0.owl (imports Core)
└── risk-ontology-extension-v1.0.owl (imports Core)
└── threatactor-ontology-extension-v1.0.owl (imports Core)
```

Critical Rule: Extensions **reference** Core, never the reverse.

7. RAG-Safe Traversal Paths

7.1 Pre-Approved Query Templates

To prevent RAG hallucination, KGCS defines **approved traversal templates**. These enforce:

1. No layer skipping
2. No circular reasoning without context
3. Only authoritative edges

Example Template T-CORE-01: "How does this CVE impact this platform?"

```
Query: CVE-X → Platform-Y?
└── CVE-X —affects—> PlatformConfiguration-Z
    (source: cvejson.configurations[].cpeMatch[].vulnerable)
└── PlatformConfiguration-Z —includes—> Platform-Y
    (source: cpematch_api_json_2.0.schema)
└── Result: Yes, if configuration matches
```

Example Template T-CORE-02: "What are all defenses against this technique?"

```

Query: Defenses for Technique-T?
└ Technique-T —mitigated_by—> DefensiveTechnique
  (source: MITRE D3FEND stix relationships)
└ Technique-T —detected_by—> DetectionAnalytic
  (source: CAR techniques[])
└ Technique-T —countered_by—> DeceptionTechnique
  (source: SHIELD mappings)

```

7.2 Forbidden Traversals

These are explicitly disallowed:

Forbidden	Why
CVE → ATT&CK (direct)	Must pass through CWE→CAPEC
Platform → CWE (direct)	Must pass through CVE
Weakness → Group (direct)	No direct relationship; contextual only
Technique → Risk Score (direct)	Risk is extension-layer only

8. Extension Layers (Incident, Risk, Threat Actor)

8.1 Why Extensions Exist

Core ontology captures **what is known to be true** (facts from standards).

Extensions capture **what we observe or infer** (facts from operations/analysis).

Critical: Extensions **never modify Core classes**. They reference them.

8.2 Incident Extension

Purpose: Track temporal, contextual observations from SIEM/SOAR

Key Classes:

- ObservedTechnique (timestamp, confidence, evidence)
- DetectionEvent (alert, rule, sensor)
- IncidentTimeline (temporal sequence)

Relationship to Core:

```

ObservedTechnique —instantiates—> Technique (ATT&CK)
DetectionEvent —detects—> ObservedTechnique
IncidentTimeline —includes—> DetectionEvent

```

Properties:

- timestamp, firstSeen, lastSeen, confidence (LOW|MEDIUM|HIGH)
- sourceSystem, evidenceIds []

8.3 Risk Extension

Purpose: Capture risk assessment, prioritization, and decision-making

Key Classes:

- RiskAssessment (scenario + score)
- RiskScenario (vulnerability → impact)
- RemediationDecision (ACCEPT|MITIGATE|TRANSFER|AVOID)

Relationship to Core:

```
RiskScenario
└── involves--> Vulnerability (CVE)
└── affects--> Asset
└── exploits--> AttackPattern (CAPEC)
└── mitigated_by--> DefensiveTechnique (D3FEND)
```

Properties:

- riskScore (0-100), likelihood, impact
- decision, decisionRationale, decisionDate
- owner, owner_contact

8.4 ThreatActor Extension

Purpose: Track attribution claims with confidence levels

Key Classes:

- AttributionClaim (claim about group responsibility)
- ThreatActorObservation (observed behavior)

Relationship to Core:

```
AttributionClaim
└── attributes_to--> Group (ATT&CK)
└── based_on--> ObservedTechnique
└── confidence--> (LOW|MEDIUM|HIGH|VERY_HIGH)
└── evidenceIds []--> ThreatActorObservation
```

Properties:

- claimId, timestamp, confidence
- sourceSystem (intel feed, SOAR, analyst), analyst_note

9. Implementation Roadmap

Phase 1: Core Standards (Complete □)

- CPE Ontology (v1.0)
- CVE Ontology (v1.0)
- CWE Ontology (v1.0)
- CAPEC Ontology (v1.0)
- ATT&CK Ontology (v1.0)
- D3FEND Ontology (v1.0)
- CAR Ontology (v1.0)
- SHIELD Ontology (v1.0)
- ENGAGE Ontology (v1.0)
- Core Ontology Extended (v1.0, imports all above)

Deliverable: 10 OWL/Turtle files with full 1:1 alignment to JSON/STIX schemas.

Phase 2: SHACL Validation (Next)

- Add representative positive/negative sample TTLs (`data/shacl-samples/`) and a small validator script (`scripts/validate_shacl.py`)
- Expand SHACL coverage to every OWL module and RAG template (partial: per-OWL identifier shapes added; additional shapes required)
- Add CI validation job (GitHub Actions) to run `pyshacl` on push/PR (validates changed OWL modules via manifest)
- Add governance artifacts: stable rule IDs, provenance URIs, and standard failure payload schema
- Integrate pre-ingest and pre-index validation into the ETL pipeline
- Add representative positive/negative sample TTLs (`data/shacl-samples/`) and a small validator script (`scripts/validate_shacl.py`)
- Expand SHACL coverage to every OWL module and RAG template (per-OWL identifier shapes added; further shapes may be added iteratively)
- Add CI validation job (GitHub Actions) to run `pyshacl` on push/PR (validates changed OWL modules via manifest)
- Add governance artifacts: stable rule IDs, provenance URIs, and standard failure payload schema (`docs/ontology/shacl/rule_catalog.json`, `failure_payload_schema.json`)
- Integrate pre-ingest and pre-index validation into the ETL pipeline (`scripts/ingest_pipeline.py` calls `run_validator()`) — validator emits per-file JSON reports to `artifacts/` and a consolidated index `artifacts/shacl-report-consolidated.json`
- A complete SHACL suite that covers Core ontology invariants and RAG traversal templates
- Each RAG template has at least one positive and one negative test case in `data/shacl-samples/`
- CI workflow runs `pyshacl` and validates changed OWL modules using the manifest mapping
- Validator supports RAG template selection and per-OWL bundle validation (`scripts/validate_shacl.py --template T1 / --owl attck-ontology-v1.0.owl`)
- Every SHACL constraint includes a stable rule identifier and an audit-friendly message
- A validation matrix of test cases (positive/negative) that exercise every shape
- CI that blocks non-conforming data on push/PR and produces machine-readable failure reports
- Documentation for running validation locally and for CI integration

Acceptance criteria:

- All Core SHACL shapes pass the provided "good" sample and fail the "bad" sample with expected violations
- Each RAG template has at least one positive and one negative test case in `data/shacl-samples/`
- CI workflow runs `pyshacl` and returns non-zero on validation failure
- Every SHACL constraint includes a stable rule identifier and an audit-friendly message
- Validator supports RAG template selection (e.g., `scripts/validate_shacl.py --template T1`) and shape-subset validation

Phase 3: Data Ingestion (Planned)

Overview

Phase 3 builds a repeatable, auditable ETL to ingest authoritative standards (NVD, MITRE), validate inputs with SHACL, preserve provenance and versioning, and write canonical KG nodes/edges into the selected graph store (Neo4j or RDF). The plan below is staged so we can deliver a minimal MVP quickly and iterate to full coverage.

Priority steps (MVP first)

1. Bootstrap infra & dependencies (1–2 days)

- Tasks: `requirements.txt`, `infra/docker-compose.yml` with Neo4j (or triple store) service, `scripts/setup_env.ps1` / `scripts/setup_env.sh` for local dev.
- Deliverable: reproducible dev environment and dependency manifest.
- Acceptance: `pip install -r requirements.txt` and `docker-compose up -d` start services.

2. CPE / CVE / CVSS ingest (3–5 days)

- Tasks: `scripts/ingest_cpe.py`, `scripts/ingest_nvd_cve.py` to parse NVD JSON (preserve CVSS versions as separate `VulnerabilityScore` nodes), map configurations → `PlatformConfiguration` → `Platform`.
- Deliverable: CVE nodes with linked CVSS nodes + platform configuration edges, provenance metadata (`source_uri`, `source_hash`, `ingest_time`).
- Acceptance: sample `data/*/samples` ingest produces nodes/edges visible in DB and includes provenance properties.

3. STIX ingestion (ATT&CK, D3FEND, SHIELD) (3–5 days)

- Tasks: `scripts/ingest_stix.py` using `stix2` to extract Techniques, Tactics, DefensiveTechnique and DeceptionTechnique objects and map to ontology classes and IDs.
- Deliverable: Technique/Tactic nodes and edges (`implements`, `belongs_to`, `mitigated_by`).

4. CAPEC / CWE / CAR ingestion (2–4 days)

- Tasks: XML/JSON parsers for CAPEC/CWE and CAR loader to create `AttackPattern`, `Weakness`, and `DetectionAnalytic` nodes and link them per the core causal chain.

5. SHACL pre-ingest and pipeline integration (2 days)

- Tasks: call existing `run_validator()` in `scripts/ingest_pipeline.py` before index/write. On failure, emit the failure payload conforming to `docs/ontology/shacl/failure_payload_schema.json` and abort the write for offending input.
- Acceptance: bad sample aborts with a valid failure payload; good sample proceeds.

6. Versioning & re-ingest safety (2–4 days)

- Tasks: add `ingest_metadata` (job id, `source_hash`, `ingest_time`), implement transactional writes and deterministic upserts, and an edge-diff rollback procedure.
- Acceptance: reingestion is idempotent and logged with provenance.

1. Tests, CI, and artifacts (2–3 days)

- Tasks: unit tests for parsers, integration tests using `data/*/samples`, GitHub Actions job to run `validate_shacl.py` and `ingest_pipeline.py --dry-run`, publish artifacts/ and optionally fail the job on configured `rule_id` values.
- Deliverable: `.github/workflows/ingest-and-validate.yml`, `tests/`, `docs/PHASE3.md` runbook.

Performance & scale (optional)

- Run batched ingest on larger samples, measure resource usage, and optimize SHACL checks or move expensive checks to offline batch jobs.

Acceptance criteria (Phase 3 complete)

- Core standards (CPE, CVE, CVSS, CWE, CAPEC, ATT&CK, D3FEND, CAR, SHIELD, ENGAGE) can be ingested into the KG with provenance recorded.
- SHACL validation runs pre-ingest and blocks invalid inputs, emitting machine-readable failure payloads.
- CI runs a dry-run ingest + SHACL validation on PRs and produces artifacts; fail-on-rule behavior is configurable.
- Re-ingestion is deterministic, auditable, and reversible via logged transactions.

Estimated timeline

- MVP (steps 1,2,5,7): 7–10 days
- Full Phase 3 (all steps): ~3–4 weeks depending on parallelization and review cadence

Deliverable: Production-ready data pipeline

Phase 4: Extension Layers (Planned)

- Incident Ontology Extension + temporal reasoning
- Risk Ontology Extension + prioritization models
- ThreatActor Ontology Extension + attribution logic
- RAG Traversal Template library

Deliverable: Safe, extensible query framework

Phase 5: AI Integration (Planned)

- RAG retrieval layer (safe traversals only)
- Explanation generation (path-based)
- Confidence scoring (traceability to source)
- Fine-tuning LLMs on KGCS queries

Deliverable: Hallucination-free cybersecurity AI

10. Why This Architecture Scales

10.1 Handling Standards Evolution

Scenario: CVSS 5.0 is released tomorrow.

Old approach: Rewrite vulnerability model, re-ingest all CVEs, hope nothing breaks.

KGCS approach:

1. Add CVSSv50Score subclass to `VulnerabilityScore` in CVE ontology
2. Load new NVD data with v5.0 metrics
3. Existing CVSS v3.1 / v4.0 nodes remain unchanged
4. RAG queries automatically see both versions

Why this works: Different CVSS versions are separate nodes with `scored_by` edges.

10.2 Adding New Standards

Scenario: You want to add NIST SP 800-53 (security controls).

KGCS approach:

1. Create `controls-ontology-v1.0.owl` (new, standalone)
2. Define edges: `DefensiveTechnique —implements—> Control`
3. Import into Core Ontology
4. Existing queries unaffected; new queries enabled

Why this works: Modular design + no inheritance = clean extension.

10.3 Supporting Organizational Context

Scenario: You have 10,000 assets to integrate.

KGCS approach:

1. Assets live in **Asset Extension**, not Core
2. Core remains pristine (no CMDB data)
3. Asset `—configured_with—> PlatformConfiguration` (from Core)
4. Asset `—affected_by—> Vulnerability` (transitive)
5. Organizational changes don't break the graph

Why this works: Clear separation of authoritative (Core) vs. contextual (Extension).

10.4 Versioning and Rollback

Scenario: You ingest bad data into CAR, want to rollback one version.

KGCS approach:

1. Every entity has `cveId`, `capecId`, `carId` (stable external ID)
2. Nodes are immutable; only edges change
3. Version control at the RDF level:

```
# Before: CAR-2013-10-002 —detects—> T1007 (with 3 evidence trails)
# Problem found: 1 trail is wrong
# After: Remove bad edge, keep 2 valid edges
# Rollback: Replay transaction, done
```

Why this works: Every relationship is traceable to source + timestamped.

11. Security Guarantees

11.1 No Hallucination

Every answer is a **path through nodes and edges**, each backed by:

- Stable external ID (CVE-YYYY-XXXX, T1234, etc.)
- Source document (JSON field, STIX property)
- Timestamp (when data was ingested)

11.2 Explainability

Every query result includes:

```
{
  "answer": "CVE-2025-1234 relates to T1059",
  "confidence": "HIGH (authoritative path)",
  "path": [
    {
      "node": "CVE-2025-1234",
      "source": "NVD JSON id field"
    },
    {
      "edge": "caused_by",
      "source": "NVD JSON weaknesses[].description[]"
    },
    ...
  ],
  "evidence_urls": ["https://nvd.nist.gov/vuln/detail/CVE-2025-1234", ...]
}
```

11.3 Auditability

- Every node change is timestamped
- Every edge change is logged with provenance
- Data can be replayed from standard sources
- No hidden inferences

11.4 Compliance

- CVSS, CWE, CAPEC, ATT&CK mappings certified by MITRE
- CPE, CVE, CVSS alignment certified by NVD
- No custom interpretations
- Suitable for:
 -

- Regulatory reporting (PCI-DSS, NIST CSF, ISO 27001)
- o SOC playbooks and runbooks
 - o Executive decision-making
-

12. Limitations and Future Work

12.1 Current Scope & Status

KGCS 1.0 covers:

- ☐ Core standards (CVE, CWE, CPE, CVSS, CAPEC, ATT&CK, D3FEND, CAR, SHIELD, ENGAGE)
- ☐ 1:1 alignment with JSON/STIX schemas
- ☐ Modular OWL ontologies
- ☐ Formal semantics and canonical invariants
- ☐ SHACL validation: core shapes, RAG shapes, AI profile, sample data, and a validator script (Phase 2 in-progress; see Section 9)
- ☐ Data ingestion pipeline (planned)
- ☐ RAG runtime enforcement & retrieval integration (planned)

12.2 Intentional Exclusions (Remains)

The following remain out of Core and belong in Extensions or operational layers:

- Threat actor motivations and attributions
- Incident timelines and raw sensor telemetry
- Asset CMDB data (organization-specific)
- Business risk scores and prioritization decisions
- SOC-specific detection rule implementations
- Exploit code and payloads (ethical/legal constraints)

12.3 Remaining Limitations and Prioritized Next Work

We have progressed SHACL validation but the project still requires targeted effort in these high-priority areas:

1. Governance & provenance: assign stable rule IDs, include provenance URLs, and standardize failure payloads for automation and audits.
2. Test matrix & coverage: add positive/negative tests for every shape and every RAG template; expand data/shacl-samples/.
3. CI integration: add a GitHub Actions job that runs `pyshacl` and publishes machine-readable reports on failure.
4. ETL integration: wire SHACL validation into pre-ingest and pre-index stages of the ingestion pipeline to prevent bad data from entering the graph.
5. Performance & scale: review and optimize SPARQL-based SHACL constraints for large graphs or move expensive checks to batch validation steps.
6. RAG runtime enforcement: implement query-time enforcement of approved traversal templates to prevent hallucinated paths.
7. Formal verification & backlog: schedule formal ontology consistency checks (reasoner runs) as part of the release process.

12.4 Long-Term Enhancements (Lower Priority)

1. Temporal reasoning and versioned edge semantics (when relationships were discovered and by whom).
2. Uncertainty quantification (explicit probability models in Risk extension).
3. Fine-grained provenance linking of CVSS metric derivations.
4. Linked-data/SPARQL publishing and subscription endpoints for downstream systems.
5. Multi-language labels and user-facing documentation improvements.

These steps complete the roadmap from "SHACL implemented" to "SHACL productionized" and will move Phase 2 from partial to done when CI, governance, and ETL integration are in place.

13. For the AI Engineer

13.1 Using KGCS for RAG

```
# Query: "What techniques can compromise this system?"
```

1. Extract asset from user context
Asset → PlatformConfiguration → CVE
2. For each CVE:
CVE → CWE → CAPEC → ATT&CK Technique
3. For each Technique:
Technique → D3FEND (defenses)
Technique → CAR (detection)
Technique → SHIELD (deception)
4. Filter by:
- Applicable to this environment (platform)
- Confidence level (HIGH for authoritative paths)
- Available data sources (CAR)
5. Return with full provenance:
[
 {
 "technique": "T1234",
 "path": [...],
 "defenses": [...],
 "detection": [...],
 "confidence": "HIGH"
 }
]

13.2 Fine-Tuning LLMs

Use KGCS to:

1. **Ground training data** in authoritative sources
2. **Filter hallucinations** (reject answers not in path)
3. **Score confidence** (authoritative path = HIGH)
4. **Generate explanations** (path = explanation)

Conclusion

KGCS is built for one purpose: Enable AI systems to reason confidently about cybersecurity.

By aligning ontologies 1:1 with standards, enforcing explicit provenance, and layering contextual extensions cleanly, KGCS ensures that:

- ☐ Every answer is traceable to source
- ☐ Standards evolve without breaking the graph
- ☐ New contexts (assets, incidents, risks) integrate cleanly
- ☐ AI can explain its reasoning with evidence
- ☐ Humans retain control and auditability

The result is an expert cybersecurity AI that is trustworthy, maintainable, and compliant.

For more information:

- See `docs/ontology/` for formal ontology specifications
- See `.github/copilot-instructions.md` for AI agent governance
- See `docs/draft/` for detailed design documents